

ASSIGNMENT 3

The Penguin Classification Analysis problem involves predicting the species of a penguin based on various physical characteristics. The dataset includes information about the body mass, culmen length, culmen depth, flipper length, and sex of different penguin species.

Clustering the data and performing classification algorithms

1. Download the dataset: Dataset
2. Load the dataset into the tool.
3. Perform Below Visualizations. • Univariate Analysis • Bi- Variate Analysis • Multi-Variate Analysis
4. Perform descriptive statistics on the dataset.
5. Check for Missing values and deal with them.
6. Find the outliers and replace them outliers
7. Check the correlation of independent variables with the target
7. Check for Categorical columns and perform encoding.
8. Split the data into dependent and independent variables.
9. Scaling the data
10. Split the data into training and testing
12. check the training and testing data shape

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

df=pd.read_csv(r"D:\MachineLearning\DataScienceCourse\
penguins_size.csv")
df
```

	species	island	culmen_length_mm	culmen_depth_mm	flipper_length_mm
0	Adelie	Torgersen	39.1	18.7	181.0
1	Adelie	Torgersen	39.5	17.4	186.0
2	Adelie	Torgersen	40.3	18.0	195.0
3	Adelie	Torgersen	NaN	NaN	NaN
4	Adelie	Torgersen	36.7	19.3	193.0
...
339	Gentoo	Biscoe	NaN	NaN	NaN

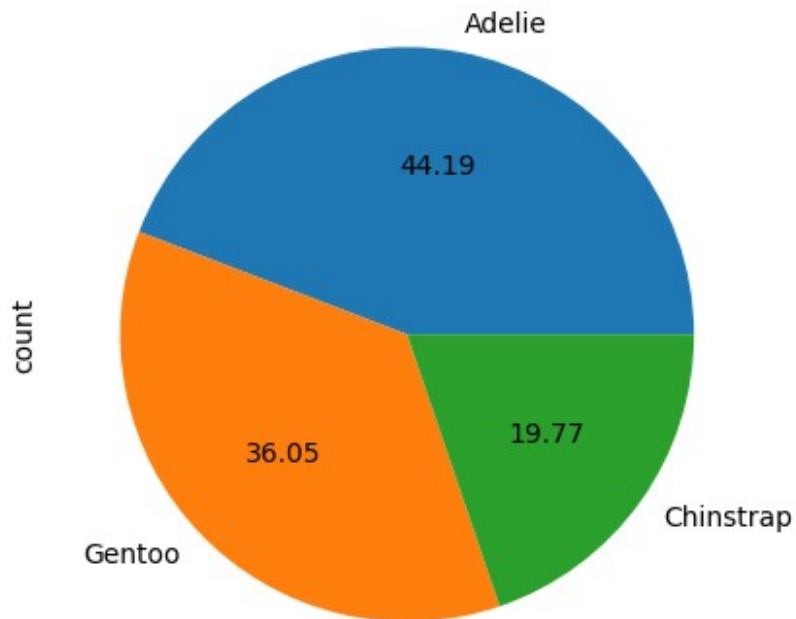
340	Gentoo	Biscoe	46.8	14.3
215.0				
341	Gentoo	Biscoe	50.4	15.7
222.0				
342	Gentoo	Biscoe	45.2	14.8
212.0				
343	Gentoo	Biscoe	49.9	16.1
213.0				

	body_mass_g	sex
0	3750.0	MALE
1	3800.0	FEMALE
2	3250.0	FEMALE
3	NaN	NaN
4	3450.0	FEMALE
...
339	NaN	NaN
340	4850.0	FEMALE
341	5750.0	MALE
342	5200.0	FEMALE
343	5400.0	MALE

[344 rows x 7 columns]

Univariate Analysis

```
df["species"].value_counts().plot(kind='pie', autopct='%.2f')
<Axes: ylabel='count'>
```



```
sns.distplot(df["culmen_length_mm"])
```

C:\Users\Vidul\AppData\Local\Temp\ipykernel_7360\2669382467.py:1:
UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df["culmen_length_mm"])
```

```
<Axes: xlabel='culmen_length_mm', ylabel='Density'>
```




Bivariate Analysis

```
sns.barplot(x=df["species"],y=df["body_mass_g"])\n<Axes: xlabel='species', ylabel='body_mass_g'>
```



```
sns.heatmap(pd.crosstab(df["species"],df["island"]))  
<Axes: xlabel='island', ylabel='species'>
```




Multivariate Analysis

```
sns.pairplot(df)
```

```
C:\Users\Vidul\AppData\Local\Programs\Python\Python311\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
```

```
self._figure.tight_layout(*args, **kwargs)
```

```
<seaborn.axisgrid.PairGrid at 0x14901835410>
```



Descriptive Statistics

```
df.describe()
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm
body_mass_g			
count	342.000000	342.000000	342.000000
mean	43.921930	17.151170	200.915205
std	5.459584	1.974793	14.061714
801.954536			

min	32.100000	13.100000	172.000000
2700.000000			
25%	39.225000	15.600000	190.000000
3550.000000			
50%	44.450000	17.300000	197.000000
4050.000000			
75%	48.500000	18.700000	213.000000
4750.000000			
max	59.600000	21.500000	231.000000
6300.000000			

Handling Missing values

```
df.isnull().sum()

species      0
island       0
culmen_length_mm    2
culmen_depth_mm    2
flipper_length_mm  2
body_mass_g      2
sex          10
dtype: int64

# Handling missing values for numerical data (using median())
l=['culmen_length_mm', 'culmen_depth_mm', 'flipper_length_mm', 'body_mass_g']
for i in l:
    df[i]=df[i].fillna(df[i].median())

# Handling missing values for categorical data (using mode)
df["sex"]=df["sex"].fillna(df["sex"].mode().iloc[0])

df.isnull().sum()
df
```

	species	island	culmen_length_mm	culmen_depth_mm
0	Adelie	Torgersen	39.10	18.7
1	Adelie	Torgersen	39.50	17.4
2	Adelie	Torgersen	40.30	18.0
3	Adelie	Torgersen	44.45	17.3
4	Adelie	Torgersen	36.70	19.3
...

339	Gentoo	Biscoe	44.45	17.3
197.0				
340	Gentoo	Biscoe	46.80	14.3
215.0				
341	Gentoo	Biscoe	50.40	15.7
222.0				
342	Gentoo	Biscoe	45.20	14.8
212.0				
343	Gentoo	Biscoe	49.90	16.1
213.0				

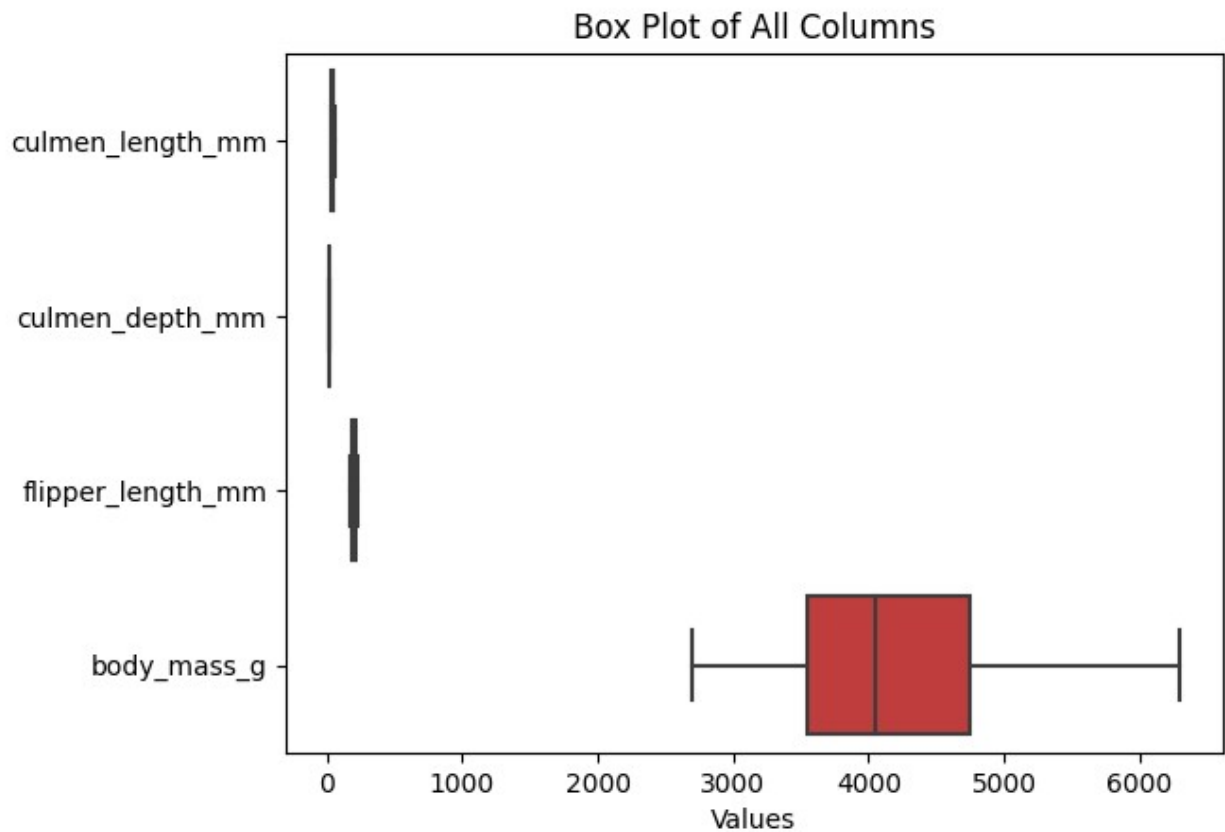
	body_mass_g	sex
0	3750.0	MALE
1	3800.0	FEMALE
2	3250.0	FEMALE
3	4050.0	MALE
4	3450.0	FEMALE
..
339	4050.0	MALE
340	4850.0	FEMALE
341	5750.0	MALE
342	5200.0	FEMALE
343	5400.0	MALE

[344 rows x 7 columns]

```
sns.boxplot(data=df, orient='h') # 'orient' is set to 'h' for
horizontal box plots
```

```
plt.xlabel('Values')
plt.title('Box Plot of All Columns')
```


```
Text(0.5, 1.0, 'Box Plot of All Columns')
```



No outliers are there

```
sns.boxplot(df["culmen_length_mm"])
```

<Axes: >



Encoding

```
# One hot
df = pd.get_dummies(df, columns = ['sex'], dtype=int)
```

```
#Label
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
l=["species","island"]
for i in l:
    df[i]=le.fit_transform(df[i])
df.head(10)
```

	species	island	culmen_length_mm	culmen_depth_mm
0	0	2	39.10	18.7
1	0	2	39.50	17.4
2	0	2	40.30	18.0
3	0	2	44.45	17.3
4	0	2	36.70	19.3
5	0	2	39.30	20.6

190.0				
6	0	2	38.90	17.8
181.0				
7	0	2	39.20	19.6
195.0				
8	0	2	34.10	18.1
193.0				
9	0	2	42.00	20.2
190.0				

	body_mass_g	sex_.	sex_FEMALE	sex_MALE
0	3750.0	0	0	1
1	3800.0	0	1	0
2	3250.0	0	1	0
3	4050.0	0	0	1
4	3450.0	0	1	0
5	3650.0	0	0	1
6	3625.0	0	1	0
7	4675.0	0	0	1
8	3475.0	0	0	1
9	4250.0	0	0	1

```
df.drop("sex_.",axis=1,inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 344 entries, 0 to 343
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	species	344 non-null	int32
1	island	344 non-null	int32
2	culmen_length_mm	344 non-null	float64
3	culmen_depth_mm	344 non-null	float64
4	flipper_length_mm	344 non-null	float64
5	body_mass_g	344 non-null	float64
6	sex_FEMALE	344 non-null	int32
7	sex_MALE	344 non-null	int32

```
dtypes: float64(4), int32(4)
```

```
memory usage: 16.3 KB
```

Checking correlations

```
df[["island","culmen_length_mm","culmen_depth_mm","flipper_length_mm",
"body_mass_g","sex_FEMALE","sex_MALE"]].corrwith(df["species"])
```

island	-0.635659
culmen_length_mm	0.728706
culmen_depth_mm	-0.741282
flipper_length_mm	0.850819

```
body_mass_g      0.747547
sex_FEMALE      -0.010240
sex_MALE         0.003185
dtype: float64
```

Splitting the data into dependent and independent variables

```
x=df.iloc[:,1:]
y=df.iloc[:,0]
x.info()
x.head(5)
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 344 entries, 0 to 343
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	island	344 non-null	int32
1	culmen_length_mm	344 non-null	float64
2	culmen_depth_mm	344 non-null	float64
3	flipper_length_mm	344 non-null	float64
4	body_mass_g	344 non-null	float64
5	sex_FEMALE	344 non-null	int32
6	sex_MALE	344 non-null	int32

```
dtypes: float64(4), int32(3)
```

```
memory usage: 14.9 KB
```

	island	culmen_length_mm	culmen_depth_mm	flipper_length_mm
0	2	39.10	18.7	181.0
				3750.0
1	2	39.50	17.4	186.0
				3800.0
2	2	40.30	18.0	195.0
				3250.0
3	2	44.45	17.3	197.0
				4050.0
4	2	36.70	19.3	193.0
				3450.0

	sex_FEMALE	sex_MALE
0	0	1
1	1	0
2	1	0
3	0	1
4	1	0

Scaling

```
# Feature scaling (MinMax Scaler or Standard Scaler)
```

```
from sklearn.pipeline import Pipeline
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler=MinMaxScaler()
```

```
x=scaler.fit_transform(x)
```

```
x=pd.DataFrame(x)
```

```
x
```

	0	1	2	3	4	5	6
0	1.0	0.254545	0.666667	0.152542	0.291667	0.0	1.0
1	1.0	0.269091	0.511905	0.237288	0.305556	1.0	0.0
2	1.0	0.298182	0.583333	0.389831	0.152778	1.0	0.0
3	1.0	0.449091	0.500000	0.423729	0.375000	0.0	1.0
4	1.0	0.167273	0.738095	0.355932	0.208333	1.0	0.0
...
339	0.0	0.449091	0.500000	0.423729	0.375000	0.0	1.0
340	0.0	0.534545	0.142857	0.728814	0.597222	1.0	0.0
341	0.0	0.665455	0.309524	0.847458	0.847222	0.0	1.0
342	0.0	0.476364	0.202381	0.677966	0.694444	1.0	0.0
343	0.0	0.647273	0.357143	0.694915	0.750000	0.0	1.0

```
[344 rows x 7 columns]
```

Train,Test,Split

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=32)
```

```
x_train.shape
```

```
(275, 7)
```

```
x_test.shape
```

```
(69, 7)
```

```
y_train.shape
```

```
(275,)
```

```
y_test.shape
```

```
(69,)
```