

Vidul Garg

## Project Title:

Grapes to Greatness: Machine Learning in Wine Quality Prediction

## Description:

Predicting wine quality using machine learning is a common and valuable application in the field of data science and analytics. Wine quality prediction involves building a model that can assess and predict the quality of a wine based on various input features, such as chemical composition, sensory characteristics, and environmental factors.

## Tasks:

Load the Dataset, Data preprocessing including visualization, Machine Learning Model building, Evaluate the model, Test with random observation

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.metrics import classification_report

df=pd.read_csv(r"D:\MachineLearning\DataScienceCourse\winequality-
red.csv")
df
```

	fixed acidity	volatile acidity	citric acid	residual sugar
0	7.4	0.700	0.00	1.9
1	7.8	0.880	0.00	2.6
2	7.8	0.760	0.04	2.3
3	11.2	0.280	0.56	1.9
4	7.4	0.700	0.00	1.9
...	...	...	...	...
...				

1594	6.2	0.600	0.08	2.0
0.090				
1595	5.9	0.550	0.10	2.2
0.062				
1596	6.3	0.510	0.13	2.3
0.076				
1597	5.9	0.645	0.12	2.0
0.075				
1598	6.0	0.310	0.47	3.6
0.067				

	free sulfur dioxide	total sulfur dioxide	density	pH
sulphates \				
0	11.0	34.0	0.99780	3.51
0.56				
1	25.0	67.0	0.99680	3.20
0.68				
2	15.0	54.0	0.99700	3.26
0.65				
3	17.0	60.0	0.99800	3.16
0.58				
4	11.0	34.0	0.99780	3.51
0.56				
...	...	...	...	...
...				
1594	32.0	44.0	0.99490	3.45
0.58				
1595	39.0	51.0	0.99512	3.52
0.76				
1596	29.0	40.0	0.99574	3.42
0.75				
1597	32.0	44.0	0.99547	3.57
0.71				
1598	18.0	42.0	0.99549	3.39
0.66				

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5
...	...	...
...		
1594	10.5	5
1595	11.2	6
1596	11.0	6
1597	10.2	5
1598	11.0	6

[1599 rows x 12 columns]

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

## Checking null values

```
df.isnull().sum()

fixed acidity          0
volatile acidity       0
citric acid            0
residual sugar         0
chlorides              0
free sulfur dioxide    0
total sulfur dioxide   0
density                0
pH                    0
sulphates              0
alcohol                0
quality                0
dtype: int64

df.describe()

      fixed acidity  volatile acidity  citric acid  residual sugar  \
count      1599.000000      1599.000000      1599.000000      1599.000000
mean         8.319637         0.527821         0.270976         2.538806
std          1.741096         0.179060         0.194801         1.409928
min          4.600000         0.120000         0.000000         0.900000
25%          7.100000         0.390000         0.090000         1.900000
50%          7.900000         0.520000         0.260000         2.200000
75%          9.200000         0.640000         0.420000         2.600000
```

max	15.900000	1.580000	1.000000	15.500000
-----	-----------	----------	----------	-----------

	chlorides	free sulfur dioxide	total sulfur dioxide
density \			
count	1599.000000	1599.000000	1599.000000
1599.000000			
mean	0.087467	15.874922	46.467792
0.996747			
std	0.047065	10.460157	32.895324
0.001887			
min	0.012000	1.000000	6.000000
0.990070			
25%	0.070000	7.000000	22.000000
0.995600			
50%	0.079000	14.000000	38.000000
0.996750			
75%	0.090000	21.000000	62.000000
0.997835			
max	0.611000	72.000000	289.000000
1.003690			


  

	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	3.311113	0.658149	10.422983	5.636023
std	0.154386	0.169507	1.065668	0.807569
min	2.740000	0.330000	8.400000	3.000000
25%	3.210000	0.550000	9.500000	5.000000
50%	3.310000	0.620000	10.200000	6.000000
75%	3.400000	0.730000	11.100000	6.000000
max	4.010000	2.000000	14.900000	8.000000

## Data Visualization

```
plt.figure(figsize=(5,3))
df["quality"].value_counts().plot(kind='bar')
plt.xticks(rotation=0)
```

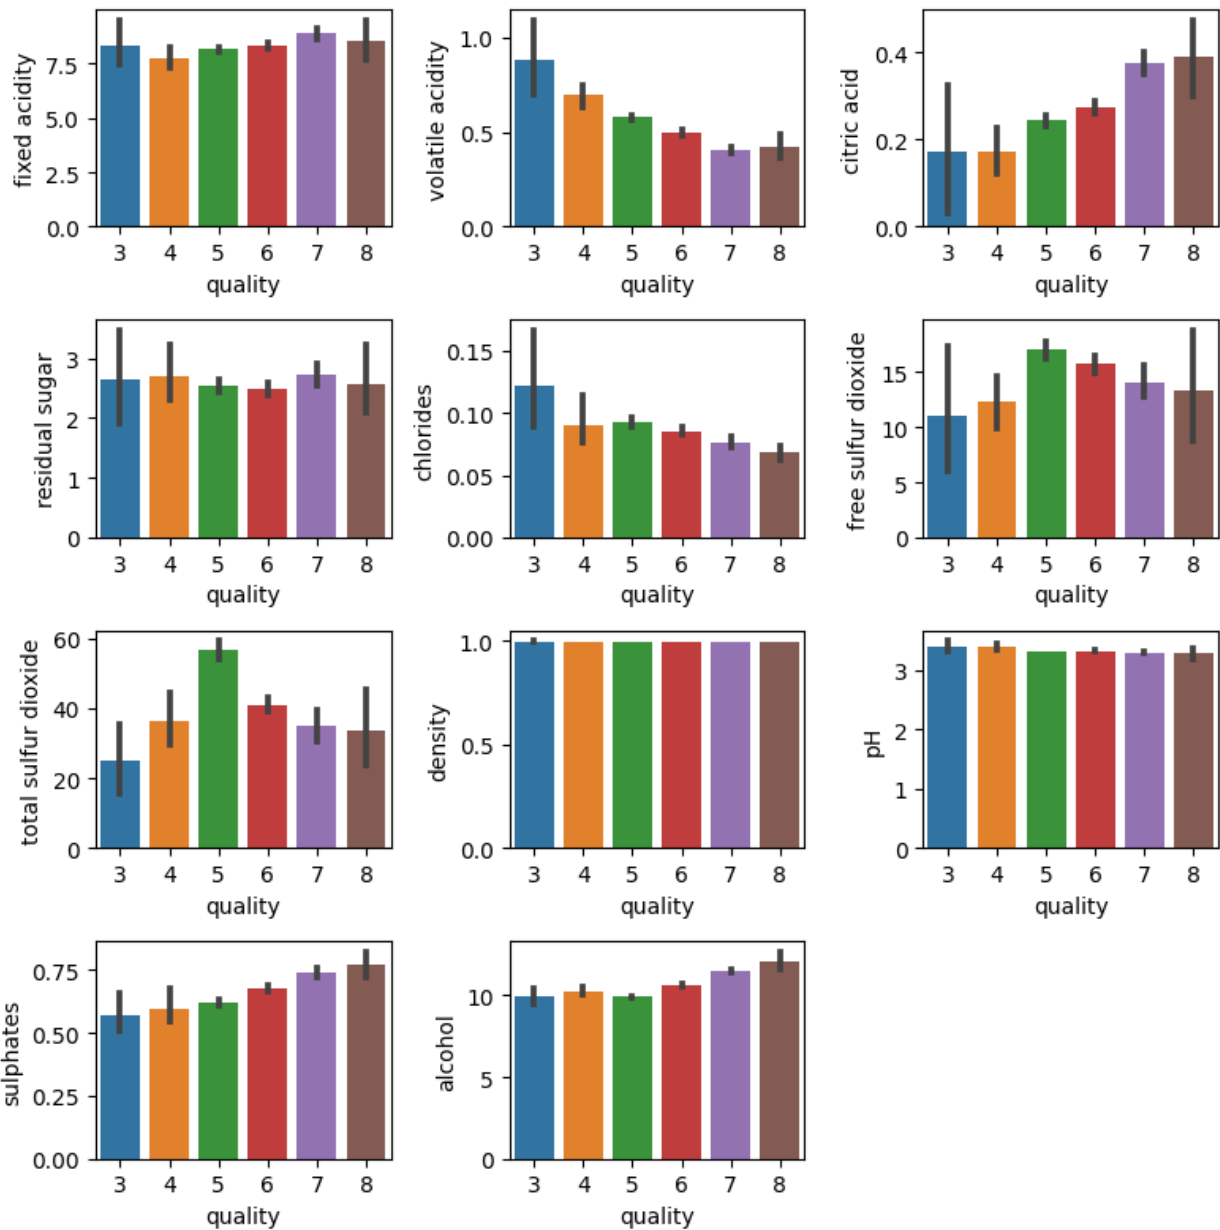
```
(array([0, 1, 2, 3, 4, 5]),
 [Text(0, 0, '5'),
  Text(1, 0, '6'),
  Text(2, 0, '7'),
  Text(3, 0, '4'),
  Text(4, 0, '8'),
  Text(5, 0, '3')])
```



Wines with quality '5' and '6' are more!!

```
plt.figure(figsize=(8,8))
l=["fixed acidity","volatile acidity","citric acid","residual
sugar","chlorides","free sulfur dioxide","total sulfur
dioxide","density","pH","sulphates","alcohol"]
for i in l:
    plt.subplot(4, 3, l.index(i) + 1)    # 4 rows, 3 columns
    sns.barplot(x=df["quality"],y=df[i])
plt.tight_layout()


# sns.barplot(x=df["quality"],y=df["alcohol"])
```



## Correlation Check

```
plt.figure(figsize=(12, 8))
cor=df.corr()
sns.heatmap(cor,annot=True)
```

<Axes: >




As we can see there is no such correlated features in the dataset

## Checking outliers

```
sns.boxplot(data=df, orient='h') # 'orient' is set to 'h' for
horizontal box plots
```

```
plt.xlabel('Values')
plt.title('Box Plot of All Columns')
```

```
Text(0.5, 1.0, 'Box Plot of All Columns')
```



There are so many outliers present in the dataset !


```
l1=["fixed acidity","volatile acidity","citric acid","residual
sugar","chlorides","free sulfur dioxide","total sulfur
dioxide","density","pH","sulphates","alcohol"]
for i in l:
    q1=df[i].quantile(0.25)
    q3=df[i].quantile(0.75)
    iqr=q3-q1
    upperL=q3+1.5*iqr
    lowerL=q1-1.5*iqr

df[i]=np.where(df[i]>upperL,upperL,np.where(df[i]<lowerL,lowerL,df[i]))

sns.boxplot(data=df, orient='h') # 'orient' is set to 'h' for
horizontal box plots

plt.xlabel('Values')
plt.title('Box Plot of All Columns')
Text(0.5, 1.0, 'Box Plot of All Columns')
```





```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1599 entries, 0 to 1598
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	fixed acidity	1599 non-null	float64
1	volatile acidity	1599 non-null	float64
2	citric acid	1599 non-null	float64
3	residual sugar	1599 non-null	float64
4	chlorides	1599 non-null	float64
5	free sulfur dioxide	1599 non-null	float64
6	total sulfur dioxide	1599 non-null	float64
7	density	1599 non-null	float64
8	pH	1599 non-null	float64
9	sulphates	1599 non-null	float64
10	alcohol	1599 non-null	float64
11	quality	1599 non-null	int64

```
dtypes: float64(11), int64(1)
```

```
memory usage: 150.0 KB
```

## Splitting the data into dependent and independent variables

```
x=df.iloc[:,11]
y=df.iloc[:,-1]
x.info()
y.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
dtypes: float64(11)
memory usage: 137.5 KB
<class 'pandas.core.series.Series'>
RangeIndex: 1599 entries, 0 to 1598
Series name: quality
Non-Null Count  Dtype
-----
1599 non-null   int64
dtypes: int64(1)
memory usage: 12.6 KB
```

## Train, Test, Split

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=32)
```

## Model Training

### KNN Classifier

```
model1=KNeighborsClassifier(n_neighbors=3)
model1.fit(x_train, y_train)
y_pred1 = model1.predict(x_test)
```

```
print(classification_report(y_test, y_pred1))
print(confusion_matrix(y_test, y_pred1))
```

	precision	recall	f1-score	support
3	0.00	0.00	0.00	1
4	0.00	0.00	0.00	8
5	0.46	0.58	0.51	120
6	0.54	0.44	0.48	146
7	0.38	0.30	0.33	40
8	0.00	0.00	0.00	5
accuracy			0.46	320
macro avg	0.23	0.22	0.22	320
weighted avg	0.47	0.46	0.46	320

```
[[ 0  0  1  0  0  0]
 [ 2  0  3  2  1  0]
 [ 0  6 70 37  7  0]
 [ 1  9 62 64 10  0]
 [ 0  0 15 13 12  0]
 [ 0  0  1  2  2  0]]
```

```
C:\Users\Vidul\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\metrics\_classification.py:1469:
UndefinedMetricWarning: Precision and F-score are ill-defined and
being set to 0.0 in labels with no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Vidul\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\metrics\_classification.py:1469:
UndefinedMetricWarning: Precision and F-score are ill-defined and
being set to 0.0 in labels with no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Vidul\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\metrics\_classification.py:1469:
UndefinedMetricWarning: Precision and F-score are ill-defined and
being set to 0.0 in labels with no predicted samples. Use
`zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

## Logistic Regression

```
model2=LogisticRegression(max_iter=5000)
model2.fit(x_train, y_train)
y_pred2 = model2.predict(x_test)
print(classification_report(y_test, y_pred2))
print(confusion_matrix(y_test, y_pred2))
```

	precision	recall	f1-score	support
3	0.00	0.00	0.00	1
4	0.00	0.00	0.00	8
5	0.60	0.77	0.67	120
6	0.55	0.56	0.56	146
7	0.41	0.17	0.25	40
8	0.00	0.00	0.00	5
accuracy			0.57	320
macro avg	0.26	0.25	0.25	320
weighted avg	0.53	0.57	0.54	320

```

[[ 0  0  1  0  0  0]
 [ 0  0  2  6  0  0]
 [ 0  0 92 28  0  0]
 [ 0  0 57 82  7  0]
 [ 0  0  2 31  7  0]
 [ 0  0  0  2  3  0]]

```

C:\Users\Vidul\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\metrics\\_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\Vidul\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\metrics\\_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\Vidul\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\metrics\\_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
warn\_prf(average, modifier, msg\_start, len(result))

## Decision Tree Classifier

```

model3=DecisionTreeClassifier()
model3.fit(x_train, y_train)
y_pred3 = model3.predict(x_test)
print(classification_report(y_test, y_pred3))
print(confusion_matrix(y_test,y_pred3))

```

	precision	recall	f1-score	support
3	0.00	0.00	0.00	1

	4	0.12	0.12	0.12	8
	5	0.62	0.75	0.68	120
	6	0.66	0.53	0.59	146
	7	0.36	0.40	0.38	40
	8	0.20	0.20	0.20	5
	accuracy			0.58	320
	macro avg	0.33	0.33	0.33	320
	weighted avg	0.59	0.58	0.58	320
	[[ 0 0 0 1 0 0]				
	[ 1 1 3 2 1 0]				
	[ 1 4 90 22 3 0]				
	[ 0 3 45 77 21 0]				
	[ 0 0 7 13 16 4]				
	[ 0 0 0 1 3 1]]				

## Accuracy Check

```
print("KNN Classifier Accuracy:", accuracy_score(y_test, y_pred1)*100)
print("Logistic Regression Accuracy:", accuracy_score(y_test,
y_pred2)*100)
print("Decision Tree Accuracy:", accuracy_score(y_test, y_pred3)*100)
```

```
KNN Classifier Accuracy: 45.625
Logistic Regression Accuracy: 56.562500000000001
Decision Tree Accuracy: 57.8125
```

## Predicting with random values

```
sample_check=[[6.5, 0.6, 0.3, 2.2, 0.07, 15.0, 40.0, 0.996, 3.4, 0.6,
9.5],
[8.0, 0.4, 0.4, 2.8, 0.085, 22.0, 55.0, 0.998, 3.2, 0.55,
11.2],
[6.8, 0.55, 0.15, 2.4, 0.075, 25.0, 62.0, 0.9962, 3.1,
0.75, 9.0],
[7.5, 0.45, 0.35, 2.5, 0.09, 30.0, 70.0, 0.9978, 3.5,
0.6, 11.5],
[7.0, 0.5, 0.2, 2.5, 0.08, 20.0, 60.0, 0.997, 3.3, 0.7,
10.0]
]

for i in sample_check:
    x=model2.predict([i])
    if(x>=6):
        print(x, "--> Good" )
    elif(x<6):
        print(x, "--> Not Good")
```

```
[5] --> Not Good
[6] --> Good
[5] --> Not Good
[6] --> Good
[5] --> Not Good
```

```
C:\Users\Vidul\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\base.py:464: UserWarning: X does not have valid
feature names, but LogisticRegression was fitted with feature names
warnings.warn(
```

```
C:\Users\Vidul\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\base.py:464: UserWarning: X does not have valid
feature names, but LogisticRegression was fitted with feature names
warnings.warn(
```

```
C:\Users\Vidul\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\base.py:464: UserWarning: X does not have valid
feature names, but LogisticRegression was fitted with feature names
warnings.warn(
```

```
C:\Users\Vidul\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\base.py:464: UserWarning: X does not have valid
feature names, but LogisticRegression was fitted with feature names
warnings.warn(
```

```
C:\Users\Vidul\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\base.py:464: UserWarning: X does not have valid
feature names, but LogisticRegression was fitted with feature names
warnings.warn(
```